

# Supporting the Fan-out Fan-in Pattern

---



**Mark Heath**

MICROSOFT MVP

@mark\_heath [www.markheath.net](http://www.markheath.net)




# Fan-out Fan-in Pattern

“Map Reduce”

Multiple  
Bitrates

“Sub-  
Orchestrations”



# Demo




## Implement “fan-out fan-in”

- Transcode activity takes bitrate input
- Orchestrator function calls activity with multiple bitrates
- Orchestrator waits for all transcodes to finish



# Getting Bitrates from Config



```
var bitrates =  
  GetBitratesFromConfig();
```

**Transcode Video  
Activity Function**

**Not deterministic!**



# Demo



## Getting list of bitrates from config

- Create an activity function to return bitrates



# Demo



## Creating sub-orchestrations

- Call an orchestrator from another orchestration
- Create a sub-orchestration for transcode step only



# Summary



## Implement “fan-out fan-in” pattern easily






- Build a list of tasks
- Wait with Task.WhenAll
- Activity to generate task input data

## Sub-orchestrations

- Extract smaller workflow sections into their own orchestrator function



# Fan-out/fan-in scenario in Durable Functions - Cloud backup example

📅 03/19/2018 • ⌚ 7 minutes to read • Contributors      all

*Fan-out/fan-in* refers to the pattern of executing multiple functions concurrently and then performing some aggregation on the results. This article explains a sample that uses [Durable Functions](#) to implement a fan-in/fan-out scenario. The sample is a durable function that backs up all or some of an app's site content into Azure Storage.

## Prerequisites

- [Install Durable Functions](#).
- Complete the [Hello Sequence](#) walkthrough.

## Scenario overview

In this sample, the functions upload all files under a specified directory recursively into blob storage. They also count the total number of bytes that were uploaded.

It's possible to write a single function that takes care of everything. The main problem you would run into is **scalability**. A single function execution can only run on a single VM, so the throughput will be limited by the throughput of that single VM. Another problem is **reliability**. If there's a failure midway through, or if the entire process takes more than 5 minutes, the backup could fail in a partially-completed state. It would then need to be restarted.






# Fan-out Fan-in File Copy Demo

Fan-out (map)

Fan-in (reduce)



Up next ...

Waiting for human  
interaction

